

Collaboration and Competition Project Report

Piero Macaluso

Introduction

The purpose of this document is to briefly present the work done in this project, along with the algorithm used and implemented, but also with a showcase of the results obtained.

This document can be divided into 4 parts:

- **Introduction:** the current section;
- **Learning Algorithm:** an overview of the approaches and methods used to solve the problem;
- **Results:** a presentation of the results reached with plots and GIFs of the best episodes;
- **Future Work:** some ideas on how to improve the actual algorithm.

Learning Algorithms

The *Tennis* environment consists of 2 agents in competition that interacts simultaneously with the environment.

The chosen algorithm was the Multi Agent Deep Deterministic Policy Gradient (MADDPG) one (Lowe et al. 2017). MADDPG is general-purpose multi-agent learning algorithm that leads to learned policies that only exploit local information at execution time (i.e. their observations). In addition, it does not assume a specific communication method between agents, and is applicable to cooperative, competitive and mixed interactions.

As presented in Figure 1, the authors of (Lowe et al. 2017) proposed an extension of actor-critic policy gradient methods where the critic is augmented with extra information about the policies of other agents, while the actor only has access to local information. After the completion of the training process, the algorithm only uses local actors at execution phase: this fact makes this technique equally applicable in cooperative and competitive settings.

In the context explained, the policies of other agents are needed to apply an update in Eq. 1. Knowing the observations and policies of other agents is not a particularly restrictive assumption. Our goal is to train agents to exhibit complex communicative behaviour in simulation, therefore this information is often available to all agents.

$$\mathcal{L}(\theta_i) = \mathbb{E}_{x,a,r,x'}[(Q_i^\mu(x, a_1, \dots, a_N) - y)^2] \quad (1)$$

$$y_t = r_t + Q_i^{\mu'}(x', a'_1, \dots, a'_N)|_{a'_j = \mu'_j(o_j)}$$

It is clear from Equation 1 that the loss is calculated starting from experiences sampled from a *replay buffer* that gathers all the past experiences of all agents. Therefore, $Q_i^\pi(x, a_1, \dots, a_N)$ is a centralized action-value function that takes as input the actions of all agents in addition to some state information x , and outputs the Q-value for agent i .




Figure 1: Overview of the multi-agent decentralized actor, centralized critic approach presented in (Lowe et al. 2017)

Hyper-parameters

The environment state and action space is very simple and it does not require convolutional neural networks. For this reason, both the actor and the critic were implemented with 2 fully connected layer with 128 neurons as hidden size. Both networks exploits *Leaky Rectified Linear Unit* (LeakyReLU) as non-linearity for each layer. The policy uses Tanh as non-linearity for the last layer, while the critic does not use non-linearity for the last layer.

There is a difference between the critic used in the DDPG framework and the ones we used in this context. This time, we put as input the concatenation of states and actions of all agents. The hyper-parameters used are presented in Table 1.

Taking into account the work done in the previous project,

Hyper-parameters	Value
Memory Buffer Size	10^6
Batch Size	256
Gamma	0.99
Tau	10^{-3}
Learning Rate	1×10^{-3}
Learning Frequency	2 (as #agent) updates per steps
Target Update Frequency	every 2 (as #agent) learning steps
Ornstein-Uhlenbeck Noise	$\theta = 0.15, \sigma = 0.20$ decaying

Table 1: Hyper-parameters used in the training process

we added to the algorithm *gradient clipping* for the critic, *weight initialization* for both networks and avoided the usage of *batch normalization*. For what concerns the noise, we decided to implement a degradation of the σ parameter until a minimum of 0.01 with a noise decay equal to 0.9995.

Results

The number of episodes to play was fixed at 2000 and the agent took just 1826 episodes to reach an average score on the last 100 episodes using the maximum reward over the 2 agents greater than 0.5. The highest average score on the last 100 episodes was equal to 1.36 reached at episode 2000.

To find the best solution, a test phase of 10 episodes was implemented and started every 50 training episodes to evaluate the results without taking into account exploration noise. As shown in Figure 2, the best test result was found at episode 2000¹ with an average of 2.29 over 10 episodes.




Figure 2: Training (blue) and Test (orange) Scores History

Future Work

The results obtained were very encouraging and positive, the agents managed to improve their behaviour with stability. Possible improvements ranges from the usage of Noisy Layers instead of Action Noise (Fortunato et al. 2017) to training agents with an ensemble of policies (Lowe et al. 2017).

¹GIF reporting a handful of seconds of the testing phase [here](#).




Figure 3: Training Scores of all 20 agents

References

- [Fortunato et al. 2017] Fortunato, M.; Azar, M. G.; Piot, B.; Menick, J.; Osband, I.; Graves, A.; Mnih, V.; Munos, R.; Hassabis, D.; Pietquin, O.; et al. 2017. Noisy networks for exploration. *arXiv preprint arXiv:1706.10295*.
- [Lowe et al. 2017] Lowe, R.; Wu, Y.; Tamar, A.; Harb, J.; Abbeel, P.; and Mordatch, I. 2017. Multi-agent actor-critic for mixed cooperative-competitive environments. *arXiv preprint arXiv:1706.02275*.